

- $K^*$  is an actual **data record** (with search key  $k$ )
- $k^*$  is of the form **( $k$ , RID)** - fixed length ( $k$ ,  $\bullet$ )
- $k^*$  is of the form **( $k$ , RID-list)** - e.g. ( $k$ , {RID11, RID12})

Notation	Meaning
$r$	relational algebra expression
$ r $	number of tuples in output of $r$
$ r $	number of pages in output of $r$
$b_p$	number of data records that can fit on a page
$b_i$	number of data entries that can fit on a page
$F$	average fanout of B <sup>+</sup> -tree index (i.e., number of pointers to child nodes)
$h$	height of B <sup>+</sup> -tree index (i.e., number of levels of internal nodes)
	$h = \lceil \log_F(\lceil \frac{ r }{b_i} \rceil) \rceil$ if format-2 index on table $R$
$B$	number of available buffer pages

### 04.1 Sorting and Selection

#### External Merge Sort

- sorted run**  $\rightarrow$  sorted data records written to a file on disk
- divide and conquer
  - create temporary file  $R_i$  for each  $B$  pages of  $R$  sorted
  - merge: use  $B - 1$  pages for input, 1 page for output
- total I/O  $= 2N(\lceil \log_{B-1}(N_0) \rceil + 1)$ 
  - $2N$  to create  $\lceil N/B \rceil$  sorted runs of  $B$  pages each
  - merging sorted runs:  $2N \times \lceil \log_{B-1} N_0 \rceil$

#### optimisation with blocked I/O

- sequential I/O - read/write in *buffer blocks* of  $b$  pages
- one block ( $b$  pages) for output, remaining blocks for input
- number of runs merged per pass,  $*F = \lfloor \frac{B}{b} \rfloor - 1$
- number of passes  $= \lceil \log_F(N_0) \rceil + 1$

#### Sorting with B<sup>+</sup>-trees

- when *sort key* is a *prefix of the index key* of the B<sup>+</sup>-tree
- sequentially scan leaf pages of B<sup>+</sup>-tree
  - for Format-2/3, use RID to retrieve data records

#### 04.2 SELECTION: $\sigma_p(R)$

- $\sigma_p(R)$ : selects rows from relation  $R$  satisfying predicate  $p$
- access path**: a way of accessing data records/entries
  - table scan**  $\rightarrow$  scan all data pages
  - index scan**  $\rightarrow$  scan index pages
  - index intersection**  $\rightarrow$  combine results from index scans
- selectivity** of an access path  $\rightarrow$  number of index & data pages retrieved to access data records/entries
  - more selective = fewer pages retrieved
- index  $I$  is a **covering index** for query  $Q \rightarrow$  if all attributes referenced in  $Q$  are part of the key of  $I$ 
  - $Q$  can be evaluated using  $I$  without any RID lookup (**index-only** plan)

#### Matching Predicates

- term**  $\rightarrow$  of form  $R.A \text{ op } c$  or  $R.A_i \text{ op } R.A_j$
- conjunct**  $\rightarrow$  one or more terms connected by  $\vee$
- CNF predicate**  $\rightarrow$  comprises one or more conjuncts connected by  $\wedge$

$\underbrace{(\text{rating} \geq 5 \vee \text{director} = \text{"Coen"})}_{\text{disjunctive conjunct}} \wedge \underbrace{(\text{year} > 2003)}_{\text{term/conjunct}} \wedge \underbrace{(\text{language} = \text{"English"})}_{\text{term/conjunct}}$
--

#### B<sup>+</sup>-tree matching predicates

- for index  $I = (K_1, K_2, \dots, K_n)$  and non-disjunctive CNF predicate  $p$ 
  - at most one* non-equality comparison operator which must be on the last attribute of the prefix ( $K_i$ )
- matching index: matching records are in contiguous pages
  - non-matching index: not contiguous  $\Rightarrow$  less efficient

#### Hash index matching predicates

- for hash index  $I = (K_1, K_2, \dots, K_n)$  and non-disjunctive CNF predicate  $p$ , all attributes must be present and all ops must be equality.

#### Primary/Covered Conjuncts

- primary conjuncts**  $\rightarrow$  subset of conjuncts that  $I$  matches
  - e.g.  $p = (\text{age} \geq 18) \wedge (\text{age} \leq 20) \wedge (\text{weight} = 65)$  for  $I = (\text{age}, \text{weight}, \text{height})$
- covered conjuncts**  $\rightarrow$  subset of conjuncts covered by  $I$ 
  - each attribute in covered conjuncts appears in key of  $I$
- primary conjuncts  $\subseteq$  covered conjuncts

#### Cost of Evaluation

let  $p'$  = primary conjuncts of  $p$ ,  $p_c$  = covered conjuncts of  $p$

#### B<sup>+</sup>-tree index evaluation of p

- navigate internal nodes to find first leaf page

$$\text{cost}_{\text{internal}} = \begin{cases} \lceil \log_F(\lceil \frac{|R|}{b_i} \rceil) \rceil & \text{if } I \text{ is a format-1 index} \\ \lceil \log_F(\lceil \frac{|R|}{b_i} \rceil) \rceil & \text{otherwise} \end{cases}$$

- scan leaf pages to access all qualifying data entries

$$\text{cost}_{\text{leaf}} = \begin{cases} \lceil \frac{||\sigma_{p'}(R)||}{b_i} \rceil & \text{if } I \text{ is a format-1 index} \\ \lceil \frac{||\sigma_{p'}(R)||}{b_i} \rceil & \text{otherwise} \end{cases}$$

- retrieve qualified data records via RID lookups

$$\text{cost}_{\text{RID}} = \begin{cases} 0 & \text{if } I \text{ is a covering format-1 index,} \\ ||\sigma_{p_c}(R)|| & \text{otherwise} \end{cases}$$

• reduce cost with **clustered** data records (sort RIDs):

$$\lceil \frac{||\sigma_{p_c}(R)||}{b_d} \rceil \leq \text{cost}_{RID} \leq \min\{||\sigma_{p_c}(R)||, |R|\}$$

#### hash index evaluation of p

- format-1:** cost to retrieve data records  $\geq \lceil \frac{||\sigma_{p'}(R)||}{b_d} \rceil$
- format-2:** cost to retrieve data entries  $\geq \lceil \frac{||\sigma_{p'}(R)||}{b_i} \rceil$

$$\text{cost to retrieve data records} = \begin{cases} 0 & \text{if } I \text{ is a covering index,} \\ ||\sigma_{p'}(R)|| & \text{otherwise} \end{cases}$$

## 05. PROJECTION $\pi_{A_1, \dots, A_m}(R)$ AND JOIN

- $\pi_L(R)$  eliminates duplicates,  $\pi_L^*(R)$  preserves duplicates

#### Sort-based approach

##### cost analysis

- extract attributes:  $|R|$  scan +  $|\pi_L^*(R)|$  output temp result
- sort records:  $2|\pi_L^*(R)|(\log_m(N_0) + 1)$
- remove duplicates:  $|\pi_L^*(R)|$  to scan records

#### optimised sort-based approach

- if  $B > \sqrt{|\pi_L^*(R)|}$ , same I/O cost as hash-based approach

- $N_0 = \lfloor \frac{|R|}{B} \rfloor \approx \sqrt{|\pi_L^*(R)|}$  initial sorted runs
- $\log_{B-1}(N_0) \approx 1$  merge passes

#### Hash-based approach

- partitioning phase**: hash each tuple  $t \in R$ 
  - $R = R_1 \cup R_2 \cup \dots \cup R_{B-1}$ 
    - for each  $R_i$  &  $R_j$ ,  $i \neq j$ ,  $\pi_L^*(R_i) \cap \pi_L^*(R_j) = \emptyset$
  - for each  $t$ : project attributes to form  $t'$ , hash  $h(t')$  to one output buffer, flush output buffer to disk when full
  - one buffer for input,  $(B - 1)$  buffers for output
- duplicate elimination** from each  $\pi_L^*(R_i)$ 
  - for each  $R_i$ : initialise in-mem hash table, hash each  $t \in R_i$  to bucket  $B_j$  with  $h' \neq h$ , insert if  $t \notin B_j$
  - write tuples in hash table to results
- I/O cost** (no partition overflow):  $|R| + 2|\pi_L^*(R)|$ 
  - partitioning cost:  $|R| + |\pi_L^*(R)|$
  - duplicate elimination cost:  $|\pi_L^*(R)|$
- partition overflow: recursively apply partitioning
  - to avoid,  $B >$  size of hash table for  $R_i = \frac{|\pi_L^*(R)|}{B-1} \times f$ 
    - approximately  $B > \sqrt{f \times |\pi_L^*(R)|}$

#### Projection using Indexes

- if index search key contains all wanted attributes *as a prefix*
  - index scan** data entries in order & eliminate duplicates

#### 05.2 JOIN $R \bowtie_R S$

$R$  = outer relation (smaller relation);  $S$  = inner relation

**!** for **format-2** index, add cost of retrieving record

#### nested loop joins

- tuple-based** nested loop join:  $|R| + ||R|| \times |S|$
- page-based** nested loop join:  $|R| + |R| \times |S|$
- block nested loop join**:  $|R| + (\lceil \frac{|R|}{B-2} \rceil \times |S|)$ .  $|R| \leq |S|$ 
  - 1 page output, 1 page input,  $(B - 2)$  pages to read  $R$
  - for each  $(B - 2)$  pages of  $R$ : for each  $P_S$  of  $S$ : check  $r, s$
- index nested loop join**:

$$|R| + ||R|| \times \left( \log_F(\lceil \frac{||S||}{b_d} \rceil) + \lceil \frac{||S||}{b_d ||\pi_{B_j}(S)||} \rceil \right)$$

- joining  $R(A, B) \bowtie_A S(A, C)$  with B+-tree index on  $S.A$
- for each tuple  $r \in R$ , use  $r$  to probe  $S$ 's index for match

#### sort-merge join

- sort  $R$  &  $S$ :  $2|R|(\log_m(N_R) + 1) + 2|S|(\log_m(N_S) + 1)$
- merge cost:  $|R| + |S|$  (worst case  $|R| + |R|| \times |S|$ )

#### optimised sort-merge join

- merge sorted runs until  $B > N(R, i) + N(S, j)$ ; then do merge and join at the same time
- I/O cost:  $3 \times (|R| + |S|)$ 
  - if  $B > \sqrt{2|S|}$ , one pass to merge initial sorted runs
  - $2(|R| + |S|)$  for initial sorted runs,  $|R| + |S|$  for merging

#### hash join

- partition  $R$  and  $S$  into  $k$  partitions on join column
  - $\pi_A(R_i) \cap \pi_B(S_j) = \emptyset \quad \forall R_i, S_j, i \neq j$

- $R = R_1 \cup R_2 \cup \dots \cup R_k, \quad t \in R_i \iff h(t.A) = i$
  - $S = S_1 \cup S_2 \cup \dots \cup S_k, \quad t \in S_i \iff h(t.B) = i$
- join corresponding partitions:  $R \bowtie_{R.A=S.B} S = (R_1 \bowtie S_1) \cup \dots \cup (R_k \bowtie S_k)$

#### Grace hash join

for *build relation*  $R$  and *probe relation*  $S$ ,

- partition**  $R$  and  $S$  into  $k$  partitions each,  $k = B - 1$
- probing phase**: hash  $r \in R_i$  with  $h'(r.A)$  to table  $T$ 
  - $\forall s \in S_i, r \in$  bucket  $h'(s.B)$ : output  $(r, s)$  if match
- I/O cost:  $3(|R| + |S|)$  (no partition overflow)
  - $B > \frac{f \times |R|}{B-1} + 2$  (input & output buffer)  $\approx B > \sqrt{f \times |R|}$ 
    - during probing,  $B >$  size of each partition +2
- partition overflow** if  $R_i$  cannot fit in memory
  - recursively apply partitioning to overflow partition

#### General join conditions

- multiple equality-join** conditions:  $(R.A = S.A) \wedge (R.B = S.B)$ 
  - index nested loop join: use index on some/all join attris
  - sort-merge join: sort on *combination* of attributes
  - other algos: no change
- inequality-join** conditions:  $(R.A < S.A)$ 
  - index nested loop join: requires B<sup>+</sup>-tree index
  - not applicable: sort-merge join (too much rewinding), hash-based joins
  - other algos: no change

## 06. Query Evaluation & Optimization

- Set operations can be implemented with joins and sort/hash
  - $R(A, B) \cap S(A, B) = \pi_{R.A, R.B} R \bowtie_p S$
  - where  $p = (R.A = S.A) \wedge (R.B = S.B)$
- Sorting for  $R \cup S$ 
  - Sort  $R$  and  $S$  using all attributes. Merge the sorted operands.
- Hashing for  $R \cup S$ 
  - Same as grace hash join,  $k$  partitions, when probing, discard if in, else add. Write to disk.
- Aggregation: maintain running info while scanning table.
  - Use index scan on covering index whenever possible.
- Group by: Partition by grp attr and run normal aggregation.
- Materialized evaluation**
  - Evaluate operator only when all operands are evaluated or materialized.
  - Intermediate results are written to disk which can be costly.
- Pipelined evaluation**
  - Execution of operators is interleaved. Pipelining may not always be possible.
  - Temporary relations are not stored on disk.
  - Output produced by operator is passed directly to parent.
- Blocking operator**  $O$  cannot start until all receive all input tuples from children operators
- Examples are external merge sort, smj, grace hash.
- Use **partial materialization** cheaper to read from temp output relation (due to very selective  $p$ )
- Iterator interface.
  - open**: initialize iterator state: resources, args.
  - getNext**: gen next tuple, return null when done
  - close**: Deallocate state information

- A SQL query has many equiv logical plans, which have many equiv physical plans
- Query optimization is about avoiding the worst plans, not picking the best
- Relational Algebra Equiv rules

- Commutativity of binary ops
  - $R \times S \equiv S \times R$
  - $R \bowtie S \equiv S \bowtie R$
- Associativity of binary ops
  - $(R \times S) \times T \equiv R \times (S \times T)$
  - $(R \bowtie S) \bowtie T \equiv R \bowtie (S \bowtie T)$
- Idempotence of unary op
  - $\pi_{L'}(\pi_L(R)) \equiv \pi_{L'}(R)$  if  $L' \subseteq L \subseteq \text{attr}(R)$
  - $\sigma_{p1}(\sigma_{p2}(R)) \equiv \sigma_{p1 \wedge p2}(R)$
- Commutating selection with projection
  - $\pi_L(\sigma_p(R)) \equiv \pi_L(\sigma_p(\pi_{L \cup \text{attr}(p)}(R)))$
- Commutating selection with binary ops
  - $\sigma_p(R \times S) \equiv \sigma_p(R) \times S$
  - $\sigma_p(R \bowtie_{p'} S) \equiv \sigma_p(R) \bowtie_{p'} S$
  - assuming  $\text{attr}(p) \subseteq \text{attr}(R)$
  - $\sigma_p(R \cup S) \equiv \sigma_p(R) \cup \sigma_p(S)$
- Commutating proj with binary ops

Let  $L = L_R \cup L_S$  where  $L_R \subseteq \text{attr}(R)$ ,  $L_S \subseteq \text{attr}(S)$

  - $\pi_L(R \times S) \equiv \pi_{L_R}(R) \times \pi_{L_S}(S)$
  - $\pi_L(R \bowtie_p S) \equiv \pi_{L_R}(R) \bowtie_p \pi_{L_S}(S)$  if  $\text{attr}(p) \cap \text{attr}(R) \subseteq L_R$  and  $\text{attr}(p) \cap \text{attr}(S) \subseteq L_S$
  - $\pi_L(R \cup S) \equiv \pi_L(R) \cup \pi_L(S)$
- Summary of RA optimization
  - Perform selection as early as possible, ideally before joins.

- Replace Cartesian Product by join whenever possible
- Project out useless attributes early
- If there are several joins, perform most restrictive join first.

- Types of Query Plan Trees
  - A query plan is **linear** if at least one operand for each join op is a base relation, else it is **bushy**
  - A linear plan is **left-deep** if every right operand is a base relation
  - A linear plan is **right-deep** if every left operand is a base relation
- Query plan enumeration, assume that optimal sol to smaller set is optimal sol to larger set. may not be true.
- DP algorithm for query enumeration: Calculate optimal plan for all single relations. Calculate optimal plan involving 2, 3 and so on relations by finding best way to join each smaller relation.
- System R optimizer
  - Enumerate only left-deep query plans. (too much to search o/w)
  - Avoid cross-product query plans.
  - Consider early selections and projections.
  - Uses  $dp(S_i, o_i)$  where  $o_i$  is null if unordered or a sequence of attr
  - May be cheaper if a SMJ is sorted on some sequence even if by itself is suboptimal.
- Cost estimation
  - Uniformity assumption: uniform distribution
  - Independence assumption: diff attr are independent
  - Inclusion assumption: For  $R \bowtie_{R.A=S.B} S$ , if  $\|\pi_A(R)\| \leq \|\pi_B(S)\|$  then  $\pi_A(R) \subseteq \pi_B(S)$
- Size estimation**
  - $|q| \approx |e| \times \prod_{i=1}^n rf(t_i)$ 
    - $rf(t_i) = \frac{||\sigma_i(e)||}{|e|}$ , reduction/selectivity factor
  - $rf(R = c) = \frac{|c|}{|R|}$  uniformity assumption
- Join selectivity**:  $rf(R.A = S.B) = \frac{||R \bowtie_{R.A=S.B} S||}{||R|| \times ||S||}$ 
  - Inclusion assumption: Assume that  $\|\pi_A(R)\| \leq \|\pi_B(S)\|$
  - By uniformity assumption, every  $R$ -tuple joins with  $\frac{||S||}{||\pi_B(S)||}$
  - $\therefore rf(R.A = S.B) \approx \frac{1}{\max(\|\pi_A(R)\|, \|\pi_B(S)\|)}$
- Using histograms. Partition domain into sub-ranges (**buckets**) and assume uniform value distribution within each bucket.
  - Equiwidth**: Each bucket has almost equal num values
  - Equidepth**: Each bucket has almost equal num tuples, sub-range of adj buckets can overlap
  - MCV**: Keep track of exact top-k common values and exclude from buckets

## 07. Transaction

- An active Xact is a Xact still in progress
- Schedule** = A list of actions from a set of Xacts where the order of the actions within each Xact is preserved
- Serial Schedule** = A schedule where the actions of Xacts are not interleaved
- We say that  $T_j$  **reads O** from  $T_i$  in a schedule S if the last write on O before  $R_j(O)$  in S is  $W_i(O)$
- We say that  $T_j$  reads  $T_i$  if  $T_j$  has read some object from  $T_i$
- We say that  $T_i$  performs the final write on  $O$  in a schedule S if the last write on O is  $W_i(O)$
- An interleaved Xact schedule is **correct** if it is "equivalent" to some serial schedule over the set of Xacts

Two schedules  $S \equiv_v S'$  (over the same set of Xacts) are **view equivalent** denoted by  $S \equiv_v S'$  if they satisfy:

- If  $T_i$  reads  $A$  from  $T_j$  in  $S$ , then  $T_i$  must also read  $A$  from  $T_j$  in  $S'$
- For each data object  $A$ , the Xact (if any) that performs the final write on  $A$  in  $S$  must also perform the final write on  $A$  in  $S'$

A schedule S is a **view serializable schedule (VS)** if S is view equivalent to some serial schedule over the same set of Xacts

We can also construct a directed VSG(S) to capture read-from and final-write relations among Txn. Each node in VSG repr a Txn with the following edges

- $(T_j, T_i)$  if  $T_i$  reads from  $T_j$
- $(T_j, T_i)$  if both  $T_i$  and  $T_j$  update the same O and  $T_i$  does final write on O
- $(T_j, T_i)$  if  $T_j$  read O from initial db and  $T_i$  update O

If VSG(S) is cyclic, then S is not VSS. If VSG(S) is acyclic, then S is VSS iff there exists a serial schedule produced from a topo ordering of VSG(S) that is view equivalent to S.

Two actions on the same **O conflict** if at least one is a W and both are from diff Xacts

- Dirty read (due to WR conflicts)
  - $T_2$  reads O modified by  $T_1$  and  $T_1$  has not yet committed.
  - $W_1(x), R_2(x)$
  - $T_2$  can see an inconsistent DB state
- Unrepeatable read (due to RW conflicts)
  - $T_2$  updates O that  $T_1$  reads and  $T_2$  commits while  $T_1$  is still in progress
  - $R_1(x), W_2(x), Commit_2, R_1(x)$
  - $T_1$  can get a different value if it reads O again
- Lost update (due to WW conflicts)

- $T_2$  overwrites O that was modified by  $T_1$  while  $T_1$  is still in progress
- $W_1(x), R_2(x), W_1(x), W_2(x)$
- $T_1$ 's update is lost
- 4. Phantom read
  - T re-executes a query on a predicate and gets a different set of results due to a recently committed T.
  - Can be prevented by **predicate locking**: Grant T an S on p, another T' request for X on p is blocked. *Also see index locking*

- **conflict equivalent** denoted by  $S \equiv_c S'$  if they order every pair of conflicting actions of two **committed** Xacts in the same way.
- **conflict serializable schedule (CS)** if it is conflict equivalent to a serial schedule over the same set of Xacts.
- **Conflict serializability graph** denoted as  $CSG(S) = (V, E)$  s.t
  - $V$  contains a node for each committed Xact in  $S$
  - $E$  contains  $T_i, T_j$  if an action in  $T_i$  precedes and conflicts with one of  $T_j$ 's action.

- **Theorem:** A schedule is **CS** iff its **CSG** is acyclic.
- **Theorem:**  $CSS \subseteq VSS \subseteq MVSS$ .
- **Note:** CS3223 uses **serializable** to mean **CS**
- A write on O by  $T_i$  is a **blind write** if  $T_i$  did not read O prior to the write
- **Theorem:** If S is **VS** and S has no blind writes, then S is also **CS**.
- **Cascading abort:** For correctness, if  $T_i$  read from  $T_j$  then  $T_i$  must abort if  $T_j$  aborts. We say that  $T_i$ 's abort is cascaded to  $T_i$ .
- **Recoverable schedule:** For every Xact T that commits in S,  $T'$  must commit after  $T'$  if T reads from  $T'$ .
- Recoverable schedules guarantee that committed Xacts will not be aborted but cascading aborts of active Xacts are possible.
- **Cascadeless schedule:** If whenever  $T_i$  reads from  $T_j$  in S,  $Commit_j$  must precede this read action.
- **Theorem:** A cascadeless schedule is recoverable.
- **Before-images:** Log the before-images of writes to undo the aborted Xacts. See Chap 10. But this does not always work.
  - $W_1(A), W_2(A), Abort_1$ . Undoing  $W_1(A)$  is incorrect.
- **Strict:** For every  $W_i(O)$  in S, O is not read or written by another Xact until  $T_i$  either aborts or commits.
  - Recovery using before-images is more efficient.
  - Concurrent executions are more restrictive.
- **Theorem:**  $Strict \subseteq Cascadeless \subseteq Recoverable$ .

## 08. Concurrency Control

**Transaction Scheduler:** For each input action (R, W, C, A) to the scheduler, either output action to S, postpone(block Xact) or reject (abort Xact)

Lock compatibility matrix		Lock Held			
Lock Requested	-	IS	IX	S	X
IS		✓	✓	✓	×
IX		✓	✓	✓	×
S		✓	×	✓	×
X		×	×	×	×

### Lock-Based concurrency control

1. If lock request not granted, T becomes **blocked**, execution is suspended and T is added to O's request queue.
2. When a lock on O is released, lock manager checks request of first T in request queue for O. If can be granted, T acquires lock on O and resumes after popped from queue.
3. When an Xact commits/aborts, all locks are released and T is removed from any request queue it's in.
4.  $S_i(O)$  : Xact  $T_i$  requests S on O
5.  $X_i(O)$  : Xact  $T_i$  requests X on O
6.  $U_i(O)$  : Xact  $T_i$  releases lock on O

### Two Phase Locking (2PL) Protocol

1. To read O, T must hold S or X on O.
2. To write O, T must hold X on O.
3. Once T releases a lock, T cannot request anymore.
4. 2PL = growing and shrinking phase.
5. **Theorem:** 2PL S are CS.

### Strict 2PL

- Same as 2PL points 1 and 2.
- T must hold onto locks until T commits or aborts.

• **Theorem:** Strict 2PL S are strict and CS.

### Deadlock Detection: Waits-for graph (WFG)

- Nodes represent active T.
- Add edge  $T_i \rightarrow T_j$  if  $T_i$  waiting for  $T_j$  to release lock.
- Remove edge when lock request is granted.

- Deadlock detected if WFG has a cycle.
- Break deadlock by aborting a T in the cycle.
- Alternative: timeout.

### Deadlock Prevention: Older T have higher priority than younger T

- Each T is assigned ts when started. Older T has a smaller ts.

- Suppose  $T_i$  requests for a lock that conflicts with a lock held by  $T_j$

Prevention Policy	$T_i$ has higher priority	$T_i$ has lower priority
Wait-die	$T_i$ waits for $T_j$	$T_j$ aborts
Wound-wait	$T_j$ aborts	$T_i$ waits for $T_j$

- Wait-die policy

- non-preemptive: only requesting T can be aborted
- younger T may abort repeatedly
- T that has all locks is never aborted
- Wound-wait policy is preemptive
- To avoid starvation, a restarted T must use original timestamp

**Lock Conversion** Increases concurrency by allowing lock conversions, previously serial only schedules can become interleaved.

- $UG_i(A) : T_i$  upgrades S on A to X.
- Blocked if another T is holding S on A.
- Allowed if  $T_i$  has not released any lock.
- $DG_i(A) : T_i$  downgrades X on A to S.
- Allowed if  $T_i$  has not modified A and  $T_i$  has not released any lock.

### Lock-based Isolation levels

Isolation Level	Dirty Read	Unrepeatable Read	Phantom Read	Degree	Isolation level	Write Locks	Read Locks	Predicate Locking
READ UNCOMMITTED	possible	possible	possible	0	Read Uncommitted	long duration	none	none
READ COMMITTED	not possible	not possible	possible	1	Read Committed	long duration	short duration	none
REPEATABLE READ	not possible	not possible	not possible	2	Repeatable Read	long duration	long duration	none
SERIALIZABLE	not possible	not possible	not possible	3	Serializable	long duration	long duration	yes

- **Short duration** lock acquired for an operation can be released after operation ends before T commits/abort
- **Long duration** lock acquired for an operation is held until T commits/abort
- **Lock Granularity:** Highest(coarsest) to lowest(finest): db, relation, page, tuple
- If T holds M on D, T implicitly holds M on granules finer than D.
- **Protocol:** Before acquiring S/X on G, acquire IS/IX on granules coarser than G top-down. Release locks bottom-up.

## 09. MVCC

- $W_i(O_i)$  : create new version of O denoted by  $O_i$ .  $O_0$  is initial version.
- $R_i(O_j)$  : read an appropriate version of O
- Read-only T not blocked by update T and vice versa
- Read-only T never aborted

- **Multiversion view equivalent (MVE)** denoted by  $S \equiv_{mv} S'$  if they have the same set of reads-from.
- $R_i(x_j)$  occurs in S if  $R_i(O_j)$  also occurs in S'
- **Monoversion schedule:** each read in S returns most recently created version. Also can be serial.
- **Multiversion view serializable s (MVSS):** a S where there exists a serial monoversion S' that is MVE to S.
- To show not MVSS, suppose there exists a S' and show that there is a cycle in the precede graph.

### • MVCC Protocol: Snapshot Isolation (SI)

- Each T has two timestamps  $start(T)$  and  $commit(T)$
- Each T sees a snapshot of DB that consists of updates by T' that committed before T starts
- T and T' are **concurrent** if they overlap
- $O_i$  is a **newer version** compared to  $O_j$  if  $commit(T_i) > commit(T_j)$
- If  $R_i(O)$  returns  $O_j$  then (either own update or latest version of O created by T that committed before  $T_i$  starts)
  - $j = i$  if  $W_i(O)$  precedes  $R_i(O)$  OR
  - $commit(T_j) < start(T_i)$
  - For every  $T_k, k \neq j$ , that has created a version  $O_k$  of O if  $commit(T_k) < start(T_i)$  then  $commit(T_k) < commit(T_j)$

- **Concurrent update property:** If multiple concurrent T updated same O, only one T allowed to commit. If not, S may not be serializable.

- **First Committer Wins (FCW):** Before committing T, if another committed concurrent T' that has updated some O that T has updated exists, abort T. Else commit T.

- **First Updater Wins (FUW):** When T needs to update O, T requests for X. If X is not held by another concurrent T'

- T is granted X on O. If O has been updated by any concurrent T'', T aborts. Otherwise T proceeds.
- Else if X is held by some concurrent T', T waits until T' aborts or commits.
  - If T' aborts, then assume T is granted X on O. If O has been updated by any concurrent T'', T aborts. Otherwise T proceeds.
  - If T' commits, then T is aborted.


- (FUW): When T commits/aborts, releases its X-lock(s).

- **Garbage collection** A version  $O_i$  may be deleted if there exists a newer version  $O_j$   $commit(T_i) < commit(T_j)$  s.t for every **active**  $T_k$  that started after  $T_i$  committed, we have  $commit(T_j) < start(T_k)$

- SI has similar performance to Read Committed but SI does not suffer from lost update or unrepeatable reads.

- But SI is vulnerable to non-serializable executions such as **write-skew anomaly** (left pic) and **read-only txn anomaly** (middle pic). Both pics are SI S that isn't MVSS


- SI also does not guarantee serializability. The rightmost picture shows the DSG for both anomalies (S1 = write-skew, S2 = read-only)



- **Serializable Snapshot Isolation (SSI):** Guarantees serializable SI schedules.
- Transactional dependencies

- $T_1 \rightarrow_{ww} T_2$ :  $T_1$  writes a version of x.  $T_2$  later writes **immediate successor** of x.
- $T_1 \rightarrow_{wr} T_2$ :  $T_1$  writes a version of x.  $T_2$  later reads this version of x.
- $T_1 \rightarrow_{rw} T_2$ :  $T_1$  reads a version of x.  $T_2$  later writes **immediate successor** of x.
- $x_j$  is the **immediate successor** of  $x_i$  if  $T_i$  commits before  $T_j$  and no txn that commits between  $T_i$ 's and  $T_j$ 's commits produces x.


- Maintain a Dependency Serializable Graph (DSG) to keep track of **rw dependencies** among concurrent T
- $DSG(S) = (V, E)$ ,  $V = \{T_1, \dots, T_k\}$  and  $E$
- $\rightarrow$  concurrent txn
- $\rightarrow$  non-concurrent txn
- If there exists a  $T_j$  is involved in two rw concurrent dependencies, abort one of  $T_i, T_j, T_k$ .
- May result in unnecessary rollbacks due to false positives of SI anomalies.
- **Theorem Non-MVSS SI Schedules:** If S is a SI schedule that is **not MVSS**, then
  - There is at least one cycle in  $DSG(S)$ , and
  - For each cycle in  $DSG(S)$ , there exists three txns  $T_i, T_j, T_k$  where



## 10. Crash Recovery

- Types of failure
  - Transaction failure: T aborts
  - System crash: loss of volatile memory contents
  - Media failure: data is lost on non-volatile storage
- **Recovery manager:** Supports Commit, Abort and Restart (abort all losers, and install updates from committed T not written to disk)
- desiderata: little overhead and recover quickly
- **Steal policy:** allow dirty pages updated by T to be replaced from buffer pool before T commits.
  - No steal: poor throughput. Steal: how atomic?
- **Force policy:** force dirty pages updated by T to be written to disk when T commits
  - poor response time but provide durability

	Force	No-force
<b>Steal</b>	undo & no redo	undo & redo
<b>No-steal</b>	no undo & no redo	no undo & redo



- **Log** Sequential file of records in **non-volatile/stable storage** (multi copies)
  - LSN: unique identifier, LSN of an earlier log record < later log record
  - type of log record (update, commit, abort)
  - identifier of Xact
  - pageID: what page is modified
  - prevLSN: LSN of the previous log record for same Xact
  - undoNextLSN: See CLR

### • Update log record

- id of page being updated, byte offset within page indicating when updated portion start, length: number of bytes for updated portion of page
- before-image and after-image of update

### • Compensation Log Record (CLR)

- When an update (described by a log record) is undone, create a CLR.
- Track action taken to undo update.
- undoNextLSN = LSN of next log record to be undone (this skips over the undone action in the LL)

- **Abort log record:** When Xact is to be aborted, start undoing this Xact.

- **End log record:** When the extra processing for aborted/committed Xact is done

- **Checkpoint log record:** See checkpointing

### • Transaction Table (TT)

- One entry for each **active Txn** which contains

- XactID
- lastLSN: LSN of most **recent** log record for this Xact
- Xact status (C = committed or U = has not committed)

### • Dirty Page Table (DPT)

- One entry for each **dirty page** in buffer pool which contains
  - pageID: page ID of dirty page
  - recLSN: LSN of the **earliest** log record for an update that dirtied page

### • Write-ahead logging (WAL) protocol

- Uncommitted update to db not flushed til log with before-image is flushed.
- Each db page contains LSN of most recent log record (**pageLSN**) that describes update to the page. Before flushing db page p to disk, ensure that all log r' up to log r corresponding to p's pageLSN has been flushed.
- **Force-at-commit protocol**
- Do not commit a Xact until after-images of all its updated records are in stable storage (db or log). Enforced by writing a **commit log record r** for Xact and flushing all log records (up to and including r) for Xact to disk.
- An Xact is considered committed if its commit log is written to stable storage
- **Simple checkpointing**
- Stop accepting any new update, commit and abort ops.
- Wait till all active update, commit and abort ops are done.
- Flush all dirty pages in buffer.
- Write a checkpoint log record containing Xact table.
- Resume accepting new update, commit and abort.
- Analysis phase finds latest checkpoint log and resumes from there, init Xact to checkpoint and DPT to empty.
- **Fuzzy checkpointing:** Used by ARIES in analysis start up.
- Write begin checkpoint log record (BCPLR).
- Write end checkpoint log record (ECPLR) containing current DPT and TT.
- Write a master record containing LSN of begin checkpoint to stable storage.
- We assume there are no log records between BCPLR and ECPLR.

### ARIES: Analysis

1. Retrieve begin checkpoint log record (BCPLR) identified by master record
2. Retrieve end checkpoint log record (ECPLR) corresponding to BCPLR
  - We assume there are no log records between BCPLR and ECPLR
3. Initialize DPT and TT using ECPLR's contents
4. Scan log record in forward direction to process each record r
5. If r is an **end log record**
  - Remove T from TT
6. Else
  - Add an entry in TT for T if T is not in TT
  - Update lastLSN of entry T to be r's LSN
  - Update status of entry T to C if r is a commit log record
7. If (r is a redable log record for page P) and (P is not in DPT) then
  - Create entry for P in DPT with pageID = P's pageID, recLSN = r's LSN
8. At the end of Analysis, 1) TT = list of all active Txn at time of crash, 2) DPT = **superset** of dirty pages at time of crash

### ARIES: Redo

1. Set **RedoLSN** = smallest recLSN among all dirty pages in DPT
2. Let r be the log record with LSN = RedoLSN
3. Scan the log in forward dir starting from r
4. If (r is a **redable** (update or CLR) log record) and (condition C is false)
  - Fetch page P associated with r (this requires a page access!!!)
  - If (P's pageLSN < r's LSN) then
    - Reapply logged action in r to P
    - update P's pageLSN = r's LSN
  - Else
    - Update P's entry in DPT: recLSN = P's pageLSN + 1
5. At the end of Redo Phase, 1) Create **end log records** for Xacts with status = C in TT and remove them from TT, 2) System is restored to state at time of crash. (Still have losers)
6. **Condition C:** (P is not in DPT) or (P's recLSN in DPT > r's LSN)

### ARIES: Undo

- Abort losers (active Xacts at time of crash) by undoing action in reverse order
- 2. Initialize L = set of lastLSNs (with status = U) from TT
- 3. Repeat until L becomes empty
  - 3.1. Delete largest lastLSN from L
  - 3.2. Let r be the corresponding log record
  - 3.3. If r is an **update log record** for T on page P
    - i. Create a CLR  $r_2$  for T:  $r_2$ 's undoNextLSN = r's prevLSN and prevLSN = T's lastLSN in TT
    - ii. Update T's entry in TT: lastLSN =  $r_2$ 's LSN
    - iii. Undo logged action on page P (WAL principle)
    - iv. update P's pageLSN =  $r_2$ 's LSN
    - v.  $Update = L - \text{and} - TT$  (r's prevLSN)
- 3.4. else if r is a **CLR** for T then
  - i.  $Update = L - \text{and} - TT$  (r's undoNextLSN)
- 3.5. else if r is an **abort log record**
  - i.  $Update = L - \text{and} - TT$  (r's prevLSN)
- 3.6. **def**  $Update = L - \text{and} - TT$  (lsn)
  - ii. else create an **end log record** for T and remove T from TT