

CS4231: 01-03

01. Mutual Exclusion

properties of a mutex algorithm

- mutual exclusion** → no more than one process in the critical section
- no starvation** → if a process wants to enter, it eventually can always enter
 - progress** → if one or more process wants to enter and if no one is in the critical section, then one of them can eventually enter the critical section
 - no starvation ⇒ progress
- proofs by contradiction
 - mutual exclusion: suppose both processes are in the critical section. WLOG, ...
 - progress: suppose both processes are waiting and consider the value of ...
 - no starvation: suppose process is waiting to enter the CS

peterson’s algorithm

- 2 processes only
 - indicate that you wantCS and set turn to the other process
 - spin until the other process does not wantCS OR turn == yours
- if a process wants to reenter after exiting the CS, it will let the other process (if exists) go first

```
RequestCS(0) {
    wantCS[0] = true;
    turn = 1;
    while (wantCS[1] && turn == 1) {};
}
ReleaseCS(0) {
    wantCS[0] = false;
}
```

lamport’s bakery algorithm

- for n processes: get a number first, then get served when all lower numbers have been served

hardware solutions

- disable interrupts (no context switching in CS)
- machine-level instructions executed atomically

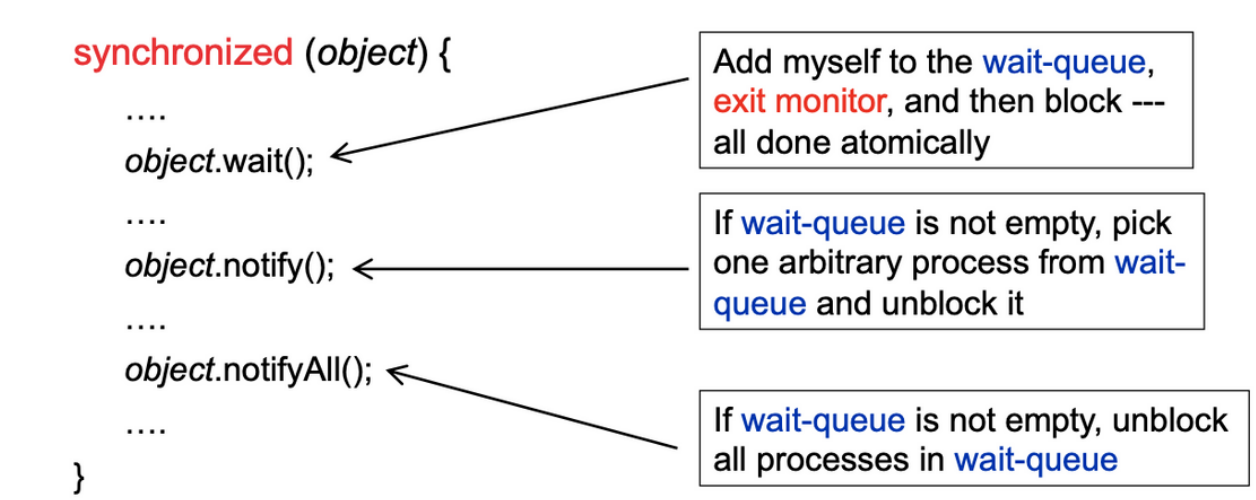
02. Synchronisation Primitives

semaphores

- no busy wait
- wait **P()** : if !value, add self to the queue and block
- signal **V()** - set value=true, wake up one arbitrary process
- avoiding deadlock
 - avoid cycles (if “wait-for” graph has a cycle, then the system has a deadlock)
 - have a total ordering

monitor

- higher-level and easier to use than semaphores
- monitor-queue:
 - enterMonitor()
 - enter monitor if no one is in
 - otherwise add self to the **monitor-queue** and block
 - exitMonitor()
 - pick one arbitrary process from **monitor-queue** and unblock
- wait-queue (within monitor):



- Hoare vs Java style monitor
 - Hoare**: notify() immediately switches from the caller to a waiting thread
 - waiting condition guaranteed to hold when waiter executes
 - Java**: notify() places a waiter on the ready thread, but signaler continues inside monitor
 - waiting condition not necessarily true when waiter executes
- caution: nested monitor

03. Consistency Conditions

definitions

- “consistent”** = satisfies the specification
- sequential consistency** → equivalent to some legal sequential history that preserves process order
 - process order** → partial order among events (same order within same process)
 - legal** → response is same as with only one process
 - sequential \nRightarrow legal
 - sequential** → invocation is *immediately* followed by response (no interleaving)
 - else *concurrent*

Sequential:
 inv(p, read, X) resp(p, read, X, 0) inv(q, write, X, 1) resp(q, write, X, OK)

concurrent:
 inv(p, read, X) inv(q, write, X, 1) resp(p, read, X, 0) resp(q, write, X, OK)
- linearisability** → sequentially consistent and preserves external order
 - external order** (<): $o1 < o2 \iff$ response of o1 is before invocation of o2
- (local property) H is linearisable $\iff \forall x, H|_x$ is linearisable
 - × sequential consistency is NOT a local property

proof

- using directed graph: a directed edge is created from o1 to o2 if
 - o1→o2 due to obj
 - o1<o2 in external order
- any topological sorting of the graph gives us a legal sequential history S
 - no cycles in the graph
- any cycle must be composed of
 - edges to some object x (\approx 1 edge since $H|x$ is equivalent S with a total order)
 - edges due to some H (\approx 1 edge since partial order induced by H is transitive)
 - edges due to some object y (or one edge)
 - edges due to some H (or one edge)

CS4231: 03-06

04. Models & Clocks

- visible orderings: **happened-before** relation (denoted $e \rightarrow f$)
 - process order, send-receive order, transitivity
- concurrent-with** relation: $e || f \iff \neg(e \rightarrow f) \wedge \neg(f \rightarrow e)$
- $e \prec f$ denotes e occurred before f in the *same process*
- $s_i \rightsquigarrow r_i$ denotes s_i is the send event corresponding to receive event r_i

logical clocks

- each process has a local counter C
 - local computation: increment C
 - send event: increment C; attach C as logical clock value V
 - receive event: $C = \max(C, V) + 1$
- $s \rightarrow t \Rightarrow C_s < C_t$ (but $C_s < C_t \nRightarrow s \rightarrow t$)

vector clock

- each event has a vector of n integers (vector clock value v)
 - $v1 = v2$ if all n fields are the same
 - $v1 \leq v2$ if all corresponding fields \leq
 - $v1 < v2$ if $v1 \leq v2$ and $v1 \neq v2$
 - $<$ is NOT a total order here!
- each process i has local vector C
 - local computation: C[i]++
 - send event: C[i]++, attach C as V
 - receive event: $C = \text{pairwise-max}(C, V)$; C[i]++
- $s \rightarrow t \iff C_s < C_t$
- proofs
 - by enumeration of all possible cases (3 happened-before cases) (same process vs diff process)
 - by contradiction

matrix clocks

- each event has n vector clocks
 - principal vector (i^{th} vector): same as vector clock
 - non-principal vector: piggybacked onto messages

05. Global Snapshot

- snapshot of local states on n process that could have happened sometime in the past
- global snapshot** \rightarrow if e2 is in the set and e1 is before e2 **in process order**, then e1 is in the set
- consistent GS** \rightarrow GS + any receive event has its corresponding send event in the set
 - can have outgoing (L \rightarrow R) arrows, but can't have incoming (R \rightarrow L) arrows
- capturing a CGS
 - special message used for snapshot protocol
 - after taking snapshot, immediately send message to other processes (total $n * (n - 1)$ messages)
- on-the-fly messages: sent before sender's local snapshot, received after receiver's local snapshot

03. Consistency Conditions (cont.)

registers: atomic \Rightarrow regular \Rightarrow safe

- atomic** register \rightarrow ensures linearisability of history
- regular** register \rightarrow when a read
 - does not overlap with any write: returns the value written by ***one of the most recent writes***
 - overlaps: returns the value written by **one of the most recent writes** OR one of the ***overlapping writes***
- safe** register \rightarrow read
 - no overlap: returns one of the most recent write values
 - overlaps: return anything
- implies nothing about sequential consistency

06. Message Ordering

- causal order** \rightarrow if s1 happened before s2, and r1 and r2 are on the same process, then r1 must be before r2
 - $s_1 \rightarrow s_2 \Rightarrow \neg(r_2 \prec r_1)$
- FIFO** \rightarrow any 2 messages from process P_i to P_j are received in the same order as they are sent
 - $s_i \prec s_j \Rightarrow \neg(r_j \prec r_i)$

protocol to maintain causal ordering

- each process maintains a $n \times n$ matrix M
 - $M[i, j]$ = # of messages sent from i to j , as known by local (current) process
 - send event (i to j): M[i, j]++; attach M as T
 - receive event:
 - if $\begin{cases} T[k, j] \leq M[k, j] & \forall k \neq i \\ T[i, j] = M[i, j] + 1 \end{cases}$, then set $M = \text{pairwise-max}(M, T)$
 - M has more knowledge about receive events at j , and exactly one message is pending from i to j
 - else, delay the message
- proofs: consider column j of matrix M and the relative values of each [k, j]

causal ordering in broadcast messages

- total ordering/atomic broadcast** \rightarrow all messages delivered to all processes in exactly the same order
- coordinator protocol**
 - one process is assigned as the coordinator
 - to broadcast: send message to coordinator; coordinator assigns a seqnum and forwards msg to all processes; messages delivered according to seqnum
 - ⊗ coordinator has too much control

skeen's algorithm for total order broadcast

- each process maintains logical clock + buffer for undelivered messages
 - process broadcasts a message
 - on receive: put the message in buffer and ACK with current logical clock value
 - sending process picks the max clock value as message number and notifies (broadcasts) message number
- deliver message if
 - all messages in the buffer have been assigned numbers
 - the message has the smallest number

proofs

- include proof that all messages will be delivered/assigned seqnums/etc